



# Software Defined Networking (SDN) in Azure Stack HCI training: Lab for Module 4: Datacenter Firewall

Microsoft Corporation  
Published: May 14, 2024

## Applies to

SDN training: Module 4: Datacenter Firewall

# Copyright

This document is provided “as-is”. Information and views expressed in this document, including URL and other Internet Web site references, may change without notice. Some examples depicted herein are provided for illustration only and are fictitious. No real association or connection is intended or should be inferred.

This document does not provide you with any legal rights to any intellectual property in any Microsoft product. You may copy, use, and modify this document for your internal reference purposes.

© 2024 Microsoft Corporation. All rights reserved.

Microsoft, Azure, Hyper-V, Internet Explorer, Silverlight, SQL Server, Windows, Windows PowerShell, and Windows Server are trademarks of the Microsoft group of companies. All other trademarks are property of their respective owners.

## Revision History

Release Date	Changes
May 14, 2024	Initial release.

# Contents

Copyright.....	2
Contents.....	3
Overview .....	4
M4.2 Datacenter firewall architecture .....	4
Lab 1: Examine firewall manager service .....	4
Lab 2: Examine NCHostAgent.....	5
M4.3 Network Security Groups .....	6
Lab 1: Create NSG for subnet .....	6
Lab 2: Create NSG for NIC .....	7
Lab 3: View ACLs (NSGs) within network controller .....	7
Lab 4: View dataplane configuration .....	7
Lab 5: Capture packet traces for firewall troubleshooting.....	9
M4.5 Trace logging.....	11
Lab 1: Enable flow audit logging.....	11

# SDN LAB: M4 DATACENTER FIREWALL

## Overview

Datacenter Firewall is a highly scalable, manageable and diagnosable software-based firewall solution that can be used to secure workloads on a virtual network. Workloads can be moved around to different hosts without breaking configuration and help secure the virtual machines, regardless of guest operating system security implementations implemented by the VM owner.

## M4.2 Datacenter firewall architecture

The Datacenter Firewall is composed of the SDN FW service within Network Controller, in addition to the FW Plugin within NCHostAgent. These services are responsible for pushing policies to dataplane to be configured and pushed into VFP where the rules are then applied to traffic traversing the port.

### Lab 1: Examine firewall manager service

1. Get the health service state by running:

```
Get-SdnServiceFabricService -ServiceTypeName 'FirewallService'
```

```
PS C:\Users\administrator> Get-SdnServiceFabricService -ServiceTypeName 'FirewallService'
```

HasPersistedState	: True
ServiceKind	: Stateful
ServiceName	: fabric:/NetworkController/FirewallService
ServiceTypeName	: FirewallService
ServiceManifestVersion	: 15.0.82
HealthState	: Ok
ServiceStatus	: Active
IsServiceGroup	: False

2. Get the process related to the Firewall Service:

```
Get-Process -Name SDNFW
```

```
[sdn-nc01]: PS C:\Users\Administrator\Documents> Get-Process -Name SDNFW
```

Handles	NPM(K)	PM(K)	WS(K)	CPU(s)	Id	SI	ProcessName
-----	-----	-----	-----	-----	--	--	-----
1222	52	82904	87912	9.73	4800	0	SDNFW

- Identify the primary replica for SDN Firewall Service:

```
Get-SdnServiceFabricReplica -ServiceTypeName 'FirewallService' -Primary
```

```
PS C:\Users\administrator> Get-SdnServiceFabricReplica -ServiceTypeName 'FirewallService' -Primary

ReplicaId           : 133132574927249843
ReplicaOrInstanceId : 133132574927249843
PartitionId         : 90db1c75-9205-4d30-8f5b-b25187fc2004
ReplicaRole         : Primary
ServiceKind         : Stateful
Id                  : 133132574927249843
ReplicaStatus       : Ready
HealthState         : Ok
ReplicaAddress      : SDN-NC03.SDN.LAB:0
NodeName            : SDN-NC03
LastInBuildDuration : 00:00:02
```

## Lab 2: Examine NCHostAgent

- Get the svchost process that NCHostAgent is running under.

```
Get-Process -Name svchost | where-Object {$_.Modules.ModuleName -icontains "nchostagent.dll"}
```

```
PS C:\Users\administrator> Get-Process -Name svchost | Where-Object {$_.Modules.ModuleName -icontains "nchostagent.dll"}

Handles  NPM(K)  PM(K)  WS(K)  CPU(s)  Id  SI ProcessName
-----  -
747      49      11692  26196   10.05   2644  0  svchost
```

- Get the TCP connections used by FW service to connect to OVSDB server. Typically, both vSwitchService and Firewall service will communicate over this port. If no workloads are present on the node, then you may only see a single TCP connection established from NC related to vSwitchService. There will not be a FirewallService connection if no workloads are deployed on the host.

```
Get-NetTCPConnection -OwningProcess 2644 -LocalPort 6640
```

```
PS C:\Users\administrator> Get-NetTCPConnection -OwningProcess 2644 -LocalPort 6640

LocalAddress      LocalPort RemoteAddress      RemotePort State      AppliedSetting OwningProcess
-----
::               6640      ::               0      Listen      Datacenter      2644
10.184.108.2     6640      10.184.108.16     53499   Established Datacenter      2644
10.184.108.2     6640      10.184.108.14     58161   Established Datacenter      2644
```

- Cross-compare the RemoteAddress and RemotePort properties to the primary replicas for vSwitchService and FirewallService to determine the correct TCP thread. For example, in [Lab 1](#), it was identified that SDN-NC03 was the primary replica for the Firewall Service.

```

PS C:\Users\Administrator> Enter-PSSession SDN-NC03
[SDN-NC03]: PS C:\Users\Administrator\Documents> Get-NetTCPConnection -LocalPort 53499

LocalAddress      LocalPort RemoteAddress      RemotePort State      AppliedSetting OwningProcess
-----
0.0.0.0           53499    0.0.0.0             0          Bound
10.184.108.16     53499    10.184.108.2        6640       Established Datacenter 428

[SDN-NC03]: PS C:\Users\Administrator\Documents> Get-Process -Id 428

Handles NPM(K) PM(K) WS(K) CPU(s) Id SI ProcessName
-----
1595    66 116432 139460 51.91 428 0 SDNFW

```

## M4.3 Network Security Groups

Network Security Groups (previously referred to as Access Control Lists) are the primary element within Network Controller that determine how the Datacenter Firewall is implemented.

### Lab 1: Create NSG for subnet

To help ensure consistent security policies for virtual machines attached to a subnet, tenant administrators may want to leverage assigning NSGs to subnets, which ensure that a consistent baseline of firewall rules are applied to all workloads associated with the subnet.

1. Navigate to **WAC** and connect to the cluster **sdnfabric.sdn.lab**.
2. Navigate to **Network Security Groups** > **+New** and specify the name: **Default**.
3. Once created, click on **Default** and under **Rules**, select **+New**.
4. Create the following rules:

Name	Priorit y	Types	Protoc ol	Source Addres s Prefix	Sourc e Port Rang e	Destinatio n Address Prefix	Destinatio n Port Range	Action s	Loggin g
TCP_80	101	Inboun d	TCP	*	*	*	80	Allow	Enable d
TCP_443	102	Inboun d	TCP	*	*	*	443	Allow	Enable d
TCP_3389	103	Inboun d	TCP	*	*	*	3389	Deny	Enable d

5. Navigate to **Virtual networks** > **Inventory** > **Contoso-VNET01** > **Subnets (Subnet01)** > **Settings**.
6. Select **Default** under **Network Security Group** and select **Submit**.

## Lab 2: Create NSG for NIC

1. Create another NSG, this time with the intent to associate to network interface of a VM directly.
2. Perform the steps 1 through 4 in [Lab 1: Create NSG for Subnet](#), with the following:
  - a. Network Security Group:
    - i. Name: Allow\_RDP
  - b. Rules:
    - i. Configure TCP\_3389 rule, with Actions = Allow (do not configure TCP\_80 / TCP\_443)
3. Navigate to **Virtual Machines > Contoso-VM2 > Settings > Networks > Network Security Group** and select **Allow\_RDP** and then **Save** your changes.
4. Attempt to perform **ping** or **Test-NetConnection -Port 3389** to the IP address of Contoso-VM1 from Contoso-VM2.
  - a. What do you observe? If you recall in Module 3 labs, ping was working between the two VMs once VNET peering was established.

## Lab 3: View ACLs (NSGs) within network controller

1. Connect to the SDN-HOST## where Contoso-VM01 is deployed.
2. Return a list of the NSGs configured within the environment.

```
$ac1s = Get-SdnResource -NcUri "https://ncnorthbound.sdn.lab" -Resource  
AccessControlLists  
$ac1s | ConvertTo-Json -Depth 10
```

- a. Locate the ACL that is connected to Contoso-VNET01/subnets/Subnet01.
- b. Locate the ACL that is connected to Contoso-VM1 NIC interface.

## Lab 4: View dataplane configuration

SDN FW service will program the underlying ACL into the MS\_Firewall database. The FW plugin within NCHostAgent will then calculate the policy and apply it to the VFP API to enforce policy on the port.

1. To confirm the ACL rules have been programmed into the dataplane correctly, navigate to SDN-HOST## where Contoso-VM1 resides.
2. Get the results of the ms\_firewall table where InstanceID in the below command is the InstanceID of the rule from previous lab of one of the rules.

```
$ovsdbAc1 = Get-SdnOvsdbFirewallRuleTable  
$ovsdbAc1 | Where-Object {$_.rule_id -ieq "InstanceID"}
```

```

vnic_id      : d1cd3748-7d2e-4c46-bd6e-ca5f773b8ac3
rule_state   : Enabled
action       : Block
dst_ports    : 3389
rule_id      : d0ea146a-e004-41d1-ade3-cc9dc40c35f3
rule_type    : RuleWithControllerPriority
direction    : Inbound
src_ports    : *
priority     : 103
logging_state : Enabled
src_ip_addresses : *
uuid         : ea32f73a-c5ea-42f9-86fe-87dce3680dad
protocols    : 6
dst_ip_addresses : *

```

- Once the rule has been confirmed to have been added into the ms\_firewall table, you can then examine the VFP rules applied to the port. We want to focus on the rules that are inbound in this scenario, since we are attempting to connect from an external endpoint into the VM.

```

$port = Get-SdnvfpVmswitchPort -VMName Contoso-VM1
Show-SdnvfpPortConfig -PortId $port.PortName -Type IPv4 -Direction IN

```

- Examine the FW\_ADMIN\_LAYER\_ID and FW\_CONTROLLER\_LAYER\_ID layers.

```

== Layer: FW_ADMIN_LAYER_ID ==
== Group: FW_GROUP_IPv4_IN_ID ==

Rule                                     Priority Flags      Type Conditions                               FlagsEx
-----
9c5dc478-252e-4be5-a12d-8f1ddf0ed0e4    103 8195 terminating stateful allow @({Protocols=6; Destination ports=3389}) 0
c5758df9-291d-482d-be4f-fe90f63822e4    65532 8195 terminating stateful block None 0
2938bdbb-d810-4175-be68-935a514476e4    65535 8195 terminating stateful allow None 0

== Layer: FW_CONTROLLER_LAYER_ID ==
== Group: FW_GROUP_IPv4_IN_ID ==

Rule                                     Priority Flags      Type Conditions                               FlagsEx
-----
d1d65ce4-30d7-4e91-a1f4-39dfb9840f54    101 8195 terminating stateful allow @({Protocols=6; Destination ports=80}) 0
69ff395b-7d07-4988-b1b0-d6fb9db61fdc    102 8195 terminating stateful allow @({Protocols=6; Destination ports=443}) 0
d0ea146a-e004-41d1-ade3-cc9dc40c35f3    103 8195 terminating stateful block @({Protocols=6; Destination ports=3389}) 0
96f55641-d9fb-4bb8-8f00-1b97961def0a    65532 8195 terminating stateful block None 0
2938bdbb-d810-4175-be68-935a514476e4    65535 8195 terminating stateful allow None 0

```

- Which layer maps to the ACLs applied to the Subnet?
  - Which layers maps to the ACLs applied to the NIC?
- To examine the VFP rules as .NET object, you can use the following:

```

# get the list of layers applied to the port
Get-SdnvfpPortLayer -PortId $port.PortName

# get the list of groups associated with a particular layer
Get-SdnvfpPortGroup -PortId $port.PortName -Layer 'FW_ADMIN_LAYER_ID'

# get the list of rules applied to a particular group
Get-SdnvfpPortRule -PortId $port.PortName -Layer 'FW_ADMIN_LAYER_ID' -Group
'FW_GROUP_IPv4_IN_ID'

```



## Lab 5: Capture packet traces for firewall troubleshooting

In some instances, viewing the configuration may not be sufficient in isolating why traffic is not routing as expected and further diagnostics are required. In these situations when troubleshooting data path issues, network traces on the vmSwitch are good next steps to further isolate what is happening within the VMSwitch and VFP layers.

1. Enable RDP on Contoso-VM1 and Contoso-VM2.

```
reg add "HKLM\SYSTEM\CurrentControlSet\Control\Terminal Server" /v  
fDenyTSConnections /t REG_DWORD /d 0 /f  
netsh advfirewall firewall set rule group="remote desktop" new enable=yes
```

2. From each of the Contoso VMs, run `Test-NetConnection -ComputerName <IPAddress> -Port 3389` against the remote VM.

```
while($true){Test-NetConnection -ComputerName 172.16.0.4 -Port 3389}  
while($true){Test-NetConnection -ComputerName 172.8.0.4 -Port 3389}
```

- a. What do you notice for behavior between the Contoso VMs that are performing the connection to 3389?
3. On SDN-HOST## where Contoso-VM1 resides performs a packet capture using the following commands, while enabling sufficient time for a repro. Alternatively, you can leverage `-ComputerName` in the cmdlet if invoking the trace remotely.

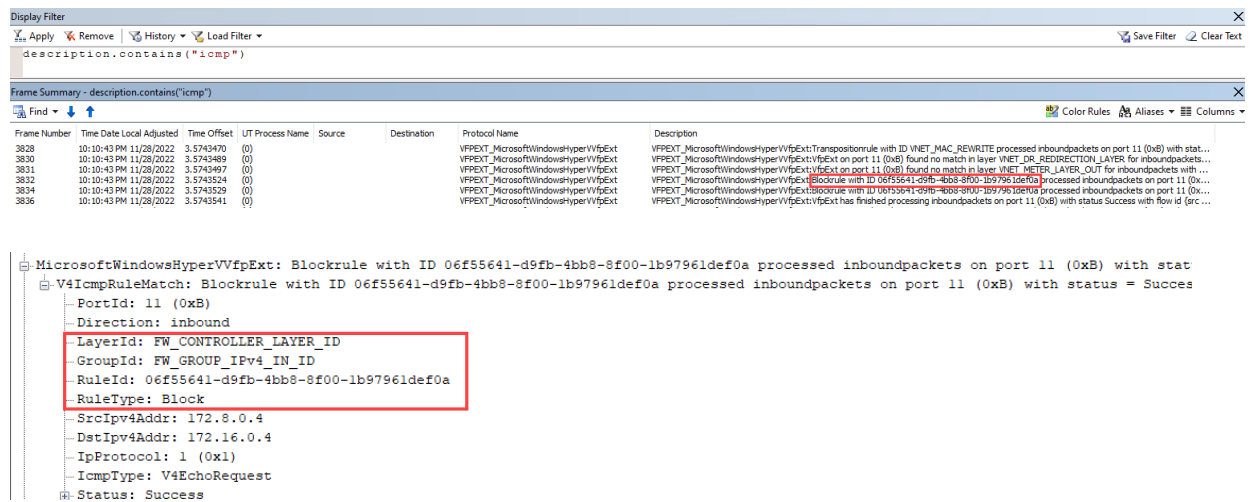
```
Start-SdnNetshTrace -Role Server
```

```
Stop-SdnNetshTrace
```

```
[sdn-host02]: PS C:\Users\administrator\Documents> Start-SdnNetshTrace -Role Server  
[SDN-HOST02] Required: 2304 MB | Available: 17978.66015625 MB  
[SDN-HOST02] Starting netsh trace  
  
Status  FileName  
-----  
Running C:\Windows\Tracing\SdnDataCollection\NetworkTraces\SDN-HOST02_20221128-225054_netshTrace.etl  
  
[sdn-host02]: PS C:\Users\administrator\Documents> Stop-SdnNetshTrace  
[SDN-HOST02] Stopping trace  
  
Status  
-----  
Stopped  
  
[sdn-host02]: PS C:\Users\administrator\Documents> _
```

4. Copy the `_netshTrace.etl` file from the host to the lab host and open it with Network Monitor 3.4.

- Similar to previously done in other labs, you will want to leverage a filter similar to `description.contains("string")` to search for patterns. In this case looking for ICMP would be sufficient.



The image shows a Wireshark packet capture. The top pane displays a filter: `description.contains("icmp")`. The middle pane shows a summary of the selected packet (Frame 3836):

- Frame Number: 3836
- Time Date Local Adjusted: 10:10:43 PM 11/28/2022
- Time Offset: 3.5743529
- UT Process Name: (0)
- Source: (0)
- Destination: (0)
- Protocol Name: VFPExt\_MicrosoftWindowsHyperVVFpExt
- Description: VFPExt\_MicrosoftWindowsHyperVVFpExt:Blockrule with ID 06f55641-d9fb-4bb8-8f00-1b97961def0a processed inboundpackets on port 11 (0xB) with status = Success

The bottom pane shows the packet details for the selected packet (Frame 3836):

- MicrosoftWindowsHyperVVFpExt: Blockrule with ID 06f55641-d9fb-4bb8-8f00-1b97961def0a processed inboundpackets on port 11 (0xB) with status = Success
- V4IcmpRuleMatch: Blockrule with ID 06f55641-d9fb-4bb8-8f00-1b97961def0a processed inboundpackets on port 11 (0xB) with status = Success
- PortId: 11 (0xB)
- Direction: inbound
- LayerId: FW\_CONTROLLER\_LAYER\_ID
- GroupId: FW\_GROUP\_IPv4\_IN\_ID
- RuleId: 06f55641-d9fb-4bb8-8f00-1b97961def0a
- RuleType: Block
- SrcIPv4Addr: 172.8.0.4
- DstIPv4Addr: 172.16.0.4
- IpProtocol: 1 (0x1)
- IcmpType: V4EchoRequest
- Status: Success

- As you can see, the packet was blocked within VFP due to rule 06f55641-d9fb-4bb8-8f00-1b97961def0a. You can cross-reference the results previously captured in previous lab, query the VFP configuration directly. Update the Layer, Group and Name parameters in the script to match your environment.

```
Get-SdnVfpPortRule -PortId $port.PortName -Layer 'FW_CONTROLLER_LAYER_ID' -
Group 'FW_GROUP_IPv4_IN_ID' -Name 06f55641-d9fb-4bb8-8f00-1b97961def0a
```

```
[sdn-host02]: PS C:\Users\administrator\Documents> Get-SdnVfpPortRule -PortId $port.PortName -Layer 'FW_CONTROLLER_LAYER_ID'
-Group 'FW_GROUP_IPv4_IN_ID' -Name 06f55641-d9fb-4bb8-8f00-1b97961def0a

Rule      : 06f55641-d9fb-4bb8-8f00-1b97961def0a
Priority   : 65532
Flags     : 8195 terminating stateful
Type      : block
Conditions : None
Flow TTL  : 240
FlagsEx   : 0
```

- Alternatively, you can use `Convert-SdnEtwTraceToText` on the host where the etl file resides to convert into .txt format, and then be able to see the same type of events as you see in Network Monitor.

```
[0]0000.0000::2022/11/28-22:10:48.752342400 [Microsoft-Windows-Hyper-V-
VfpExt]2(Block)rule with ID 06f55641-d9fb-4bb8-8f00-1b97961def0a processed
1(inbound)packets on port 11 (Name = 9811BF39-D89E-49E6-A00C-BC845608D400,
FriendlyName = NULL) with status = 0x10003(NT=??) and statusLocation = 0:
flow id {src ip = 172.8.0.4, dst ip = 172.16.0.4, protocol = 1, icmp type =
8(V4EchoRequest)}, rule {layer = FW_CONTROLLER_LAYER_ID, group =
FW_GROUP_IPv4_IN_ID, rule id = 06f55641-d9fb-4bb8-8f00-1b97961def0a,
gftFlags = 0}
```

# M4.5 Trace logging

## Lab 1: Enable flow audit logging

1. Navigate to WAC > sdnfabric.sdn.lab > Network Security Groups.
2. Select Flow Logs > Auditing Settings.
3. Configure the SDN NSG Logging Folder to C:\Windows\Tracing\Audit\_Logs and then select Save.
4. Enable Logging.

### General

SDN ACL logging folder

C:\Windows\Tracing\Audit\_Logs

Save

### SDN ACL logging usage

Host name	Log folder	Log folder usage	Disk available capacity	ACL logging status		
<a href="#">SDN-HOST02.SDN.LAB</a>	<a href="#">C:\Windows\Tracing\Audit_Logs</a>	0 B	17.54 GB	Enabled	<a href="#">Disable ACL logging</a>	<a href="#">Clean up</a>
<a href="#">SDN-HOST01.SDN.LAB</a>	<a href="#">C:\Windows\Tracing\Audit_Logs</a>	0 B	17.06 GB	Disabled	<a href="#">Enable ACL logging</a>	<a href="#">Clean up</a>

5. Confirm you see .json file in the folder on the node(s) tracing was enabled for.
  - a. Should show as 0 KB for the current active audit file. Data will be committed to the audit log every hour.